

Assumptions & Limits

What's mocked now and how a pilot would close gaps

Assumptions (MVP)

- Signals come from public sources or manually provided links (no closed APIs in MVP).
- Operators can self-identify affected assets (checkbox 'Is this us?').
- Human Impact Index weights are initial heuristics tuned with Bio-ISAC feedback.
- Tier-2 response is available for high-risk cases (Whitney/partner analysts).

Constraints / Known Gaps

- No automated asset inventory correlation in MVP; mapping is manual.
- Limited identity verification; prototype uses referral/org email checks.
- AI summaries may omit edge details; humans review before action.
- No direct EHR/LIS/SCADA integration in MVP; actions are guidance/runbook only.

Pilot Close-Out Plan

- Add asset inventory import (CSV/API) to improve 'Is this us?' accuracy.
- Implement lightweight verification (email domain + invite code + optional backchannel).
- Introduce feed automation and dedupe store; export audit logs to SIEM.
- Run A/B on Quiet Mode policy and re-weight Human Impact Index from real incidents.

Non-Goals (Phase 1)

- No automatic device shutdown controls.
- No collection of PHI/PII beyond minimal contact and org identity for verification.
- No broad ingestion of closed, confidential datasets without explicit agreements.